

# Fundamentals of Machine Learning

## Optimization

Workshop Lead: Tugce Gurbuz

Jul 24th, 2024

## What is optimization?

What do we optimize?

How do we optimize?

## What is optimization?

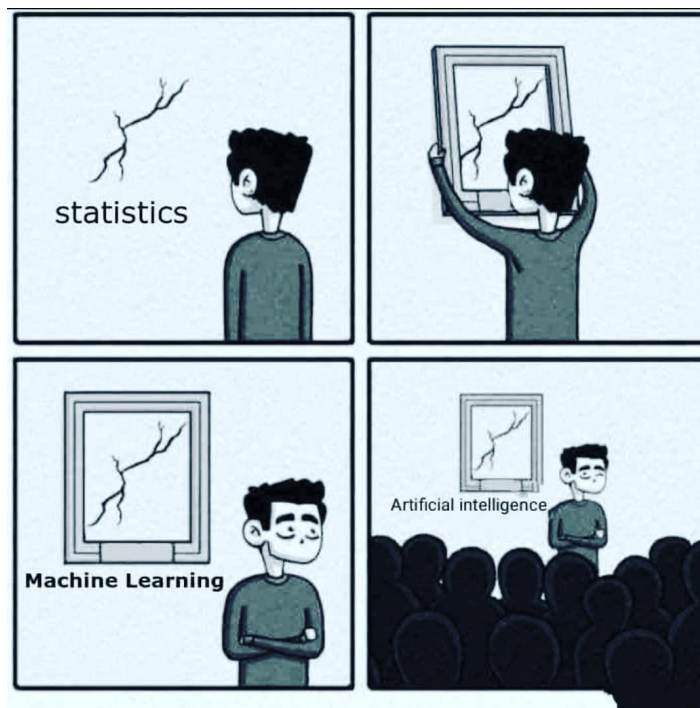
What do we optimize? -> **Parameters of the model to make the loss minimum**

How do we optimize?


## What is optimization?

What do we optimize? -> **Parameters of the model to make the loss minimum**

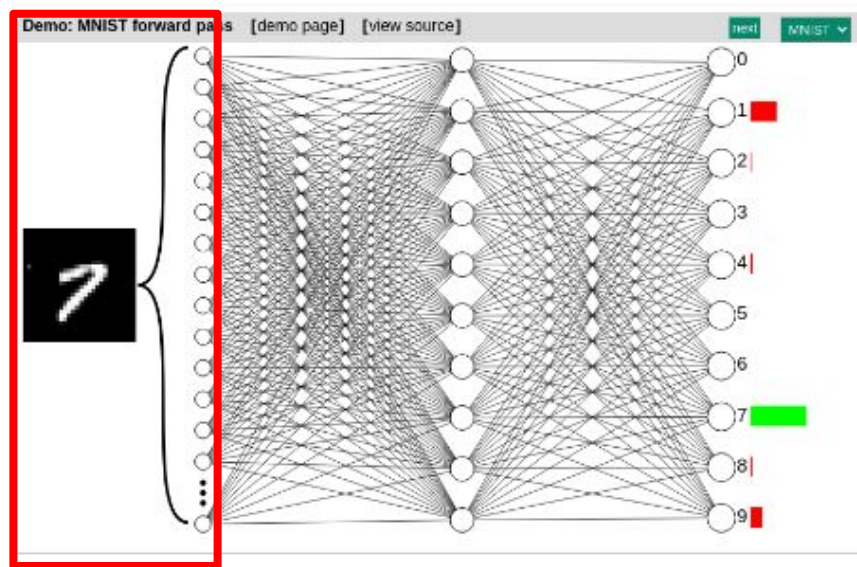
How do we optimize? -> **Math!**



## Task in tutorial: Classifying handwritten digits with MLP

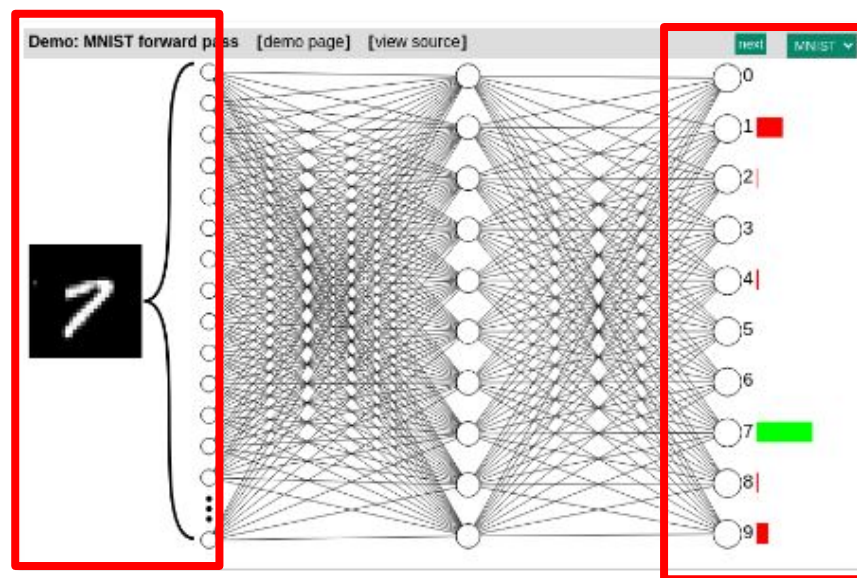


## Task in tutorial: Classifying handwritten digits with MLP



28 x 28 image -> 784 vector

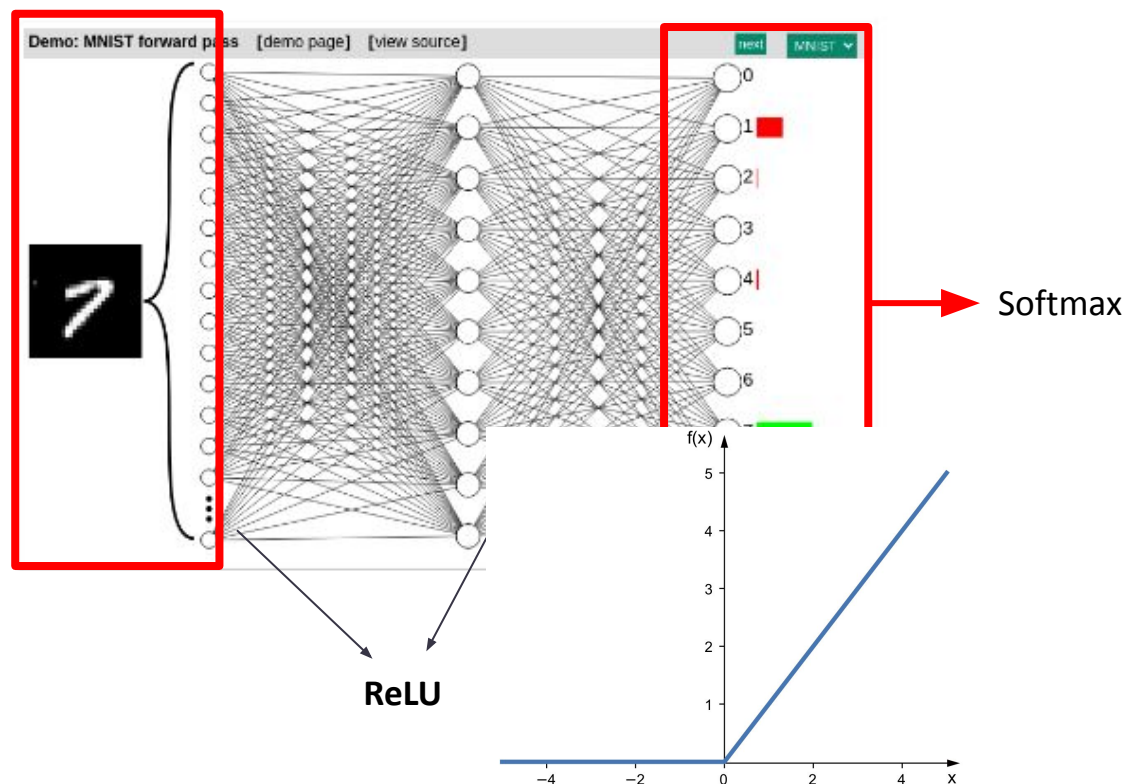
## Task in tutorial: Classifying handwritten digits with MLP



Softmax

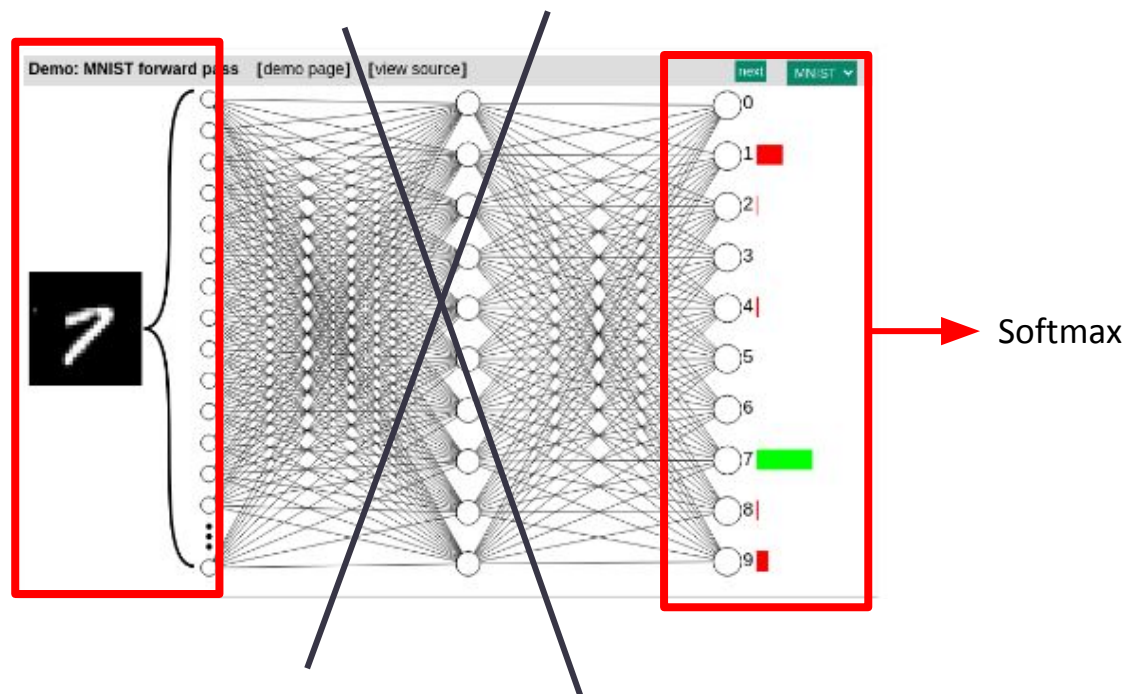
$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

## Task in tutorial: Classifying handwritten digits with MLP






## Task in tutorial: Classifying handwritten digits with MLP



## Task in tutorial: Classifying handwritten digits with MLP



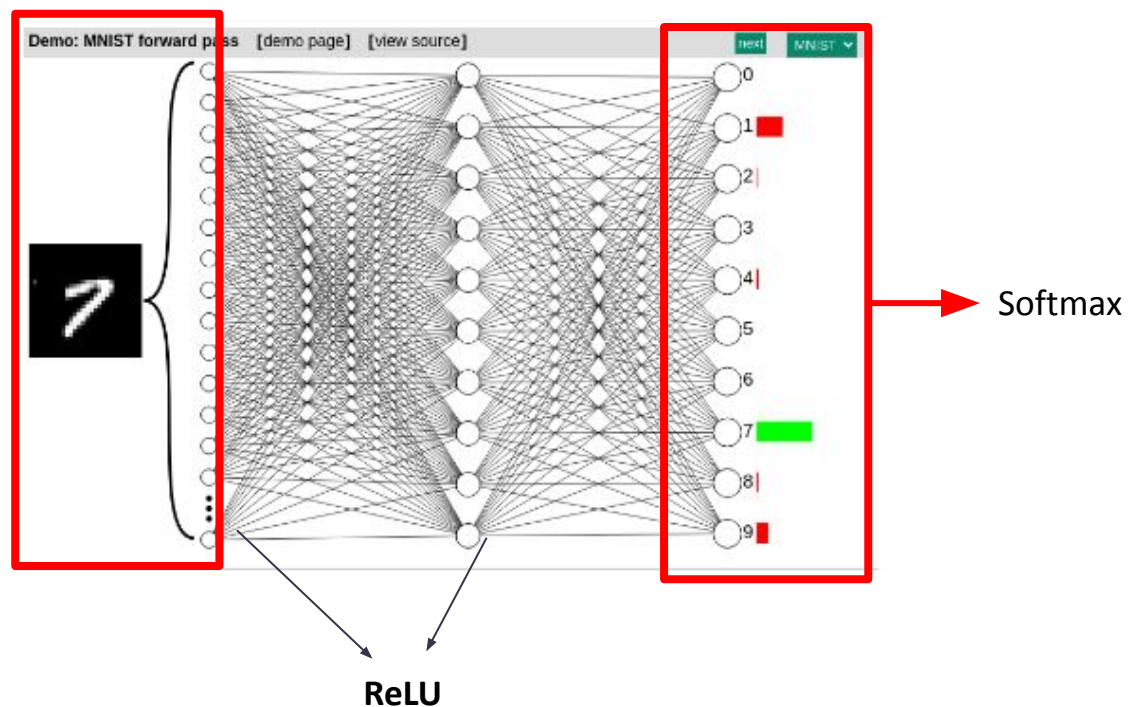
$$f(x) = \text{softmax}(Wx + b)$$

$$x \in \mathbb{R}^{784},$$

$$W \in \mathbb{R}^{10 \times 784}$$


$$b \in \mathbb{R}^{10}$$

## Task in tutorial: Classifying handwritten digits with MLP



Let's go to tutorial to implement our classification model!

## How to calculate the loss?

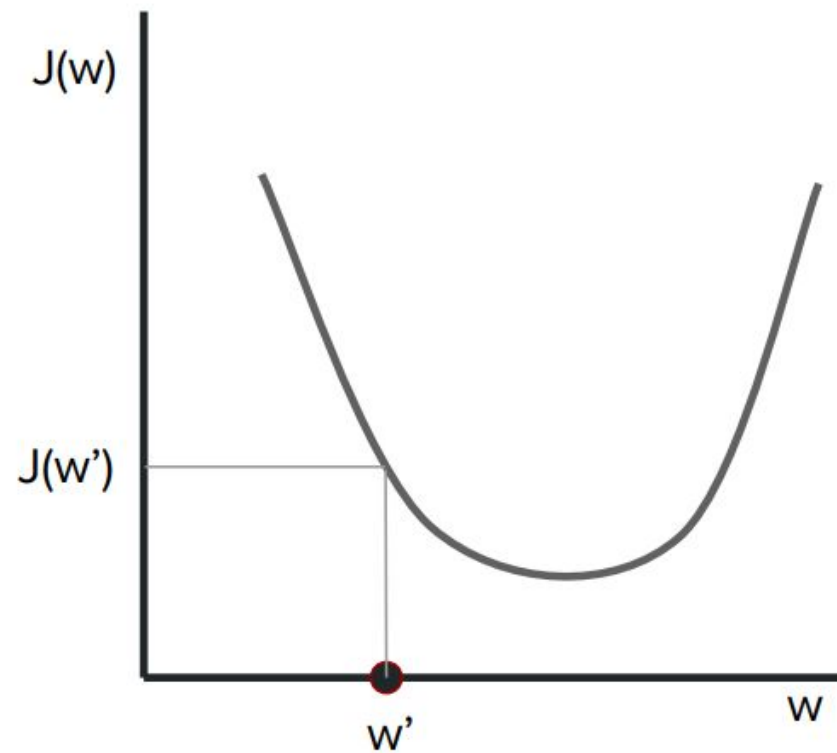


**Multiclass classification: Cross-entropy vs.  
MSE?**

## How to decide the loss function?

- Identify your problem (e.g., classification, regression?)
- Check the literature

## How to minimize the loss?



## How to minimize the loss?

Random search?

Algorithm:

sample random points around current  $w$

if random point,  $w'$ , yields lower objective (i.e.  $J(w') < J(w)$ ):

Accept  $w'$  as new position and store it in  $w$



## How to minimize the loss?

Random search?





## How to minimize the loss?

Gradient Descent <3

Algorithm:

- Compute gradient (it points uphill)
- Do step in opposite direction of gradient
- Step size (learning rate),  $\eta$



## How to minimize the loss?

Gradient Descent

Algorithm:

- Compute gradient (it points uphill)
- Do step in opposite direction of gradient
- Step size (learning rate),  $\eta$

$$w_{t+1} = w_t - \eta \nabla J(w_t)$$

## How to minimize the loss?

Gradient Descent

Algorithm:

- Compute gradient (it points uphill)
- Do step in opposite direction of gradient
- Step size (learning rate),  $\eta$

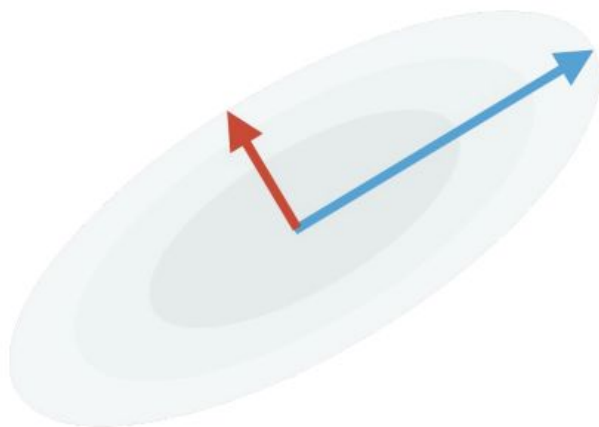
$$w_{t+1} = w_t - \eta \nabla J(w_t)$$

Let's go to our tutorial to implement Gradient Descent!!

## How to minimize the loss?

Gradient Descent

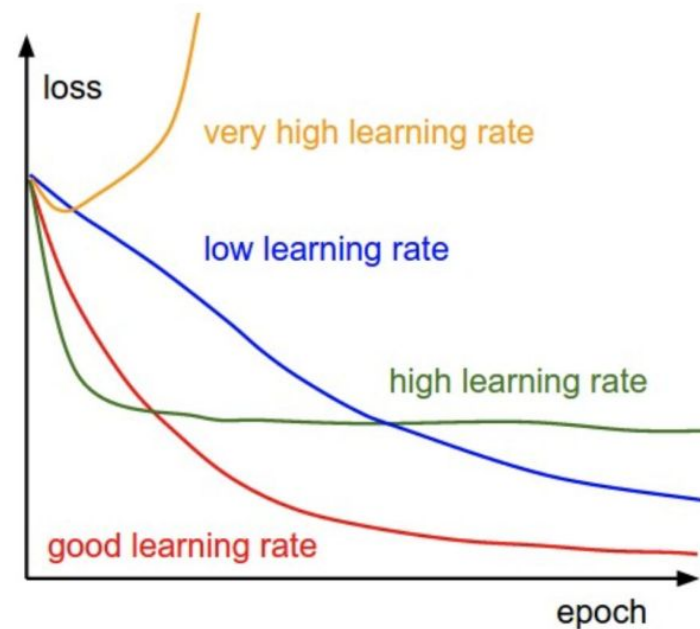
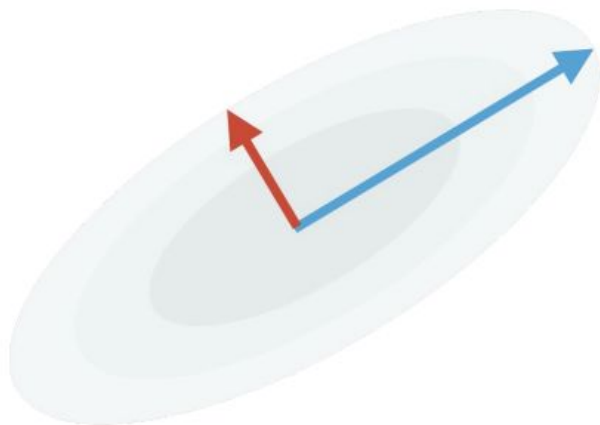
- How to choose learning rate?



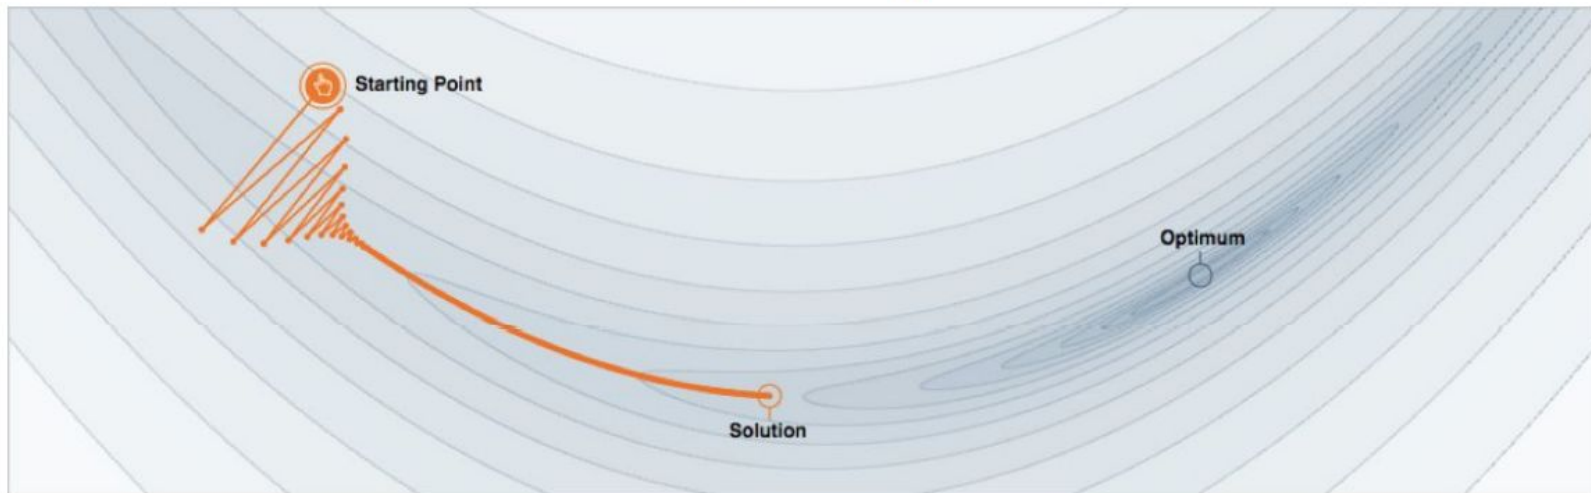
## How to minimize the loss?

Gradient Descent

- How to choose learning rate?



## How to minimize the loss?

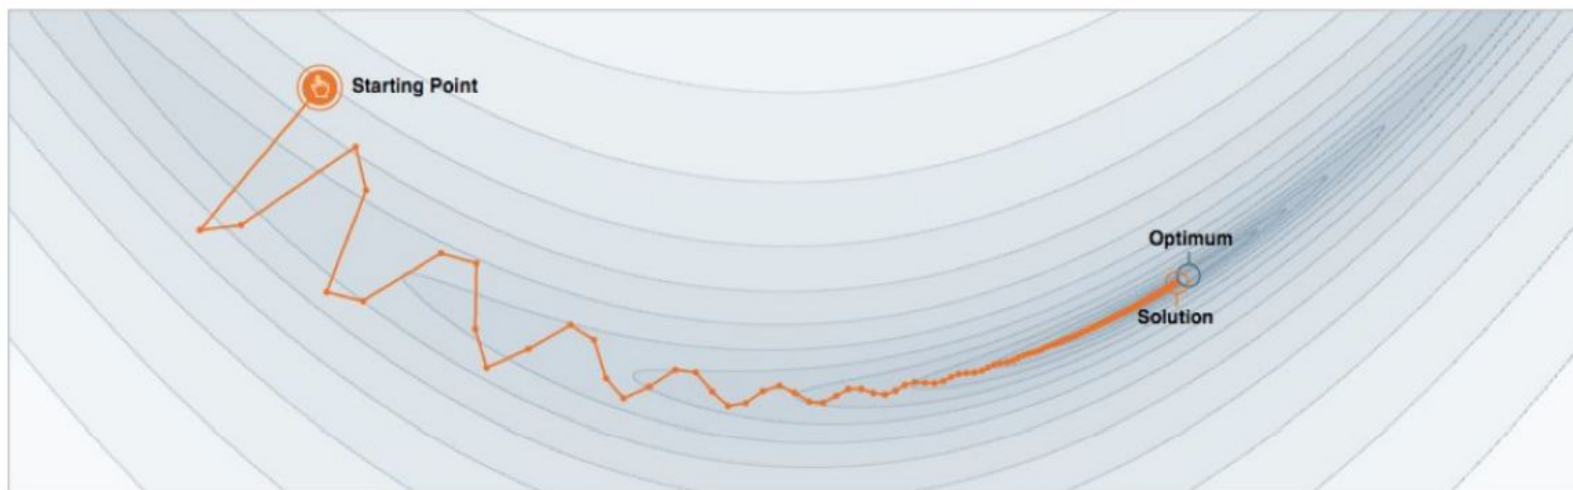


[Distill.pub]

## How to minimize the loss?

Momentum < 3

- Accelerates along flat directions
- Slows down along sharp directions



[Distill.pub]

## How to minimize the loss?

How does the momentum work?

Momentum algorithm:

- Do a gradient descent step
- Apply the update from the last iteration, only smaller (momentum step)

$$w_{t+1} = w_t - \eta \nabla J(w_t) + \beta (w_t - w_{t-1})$$



## How to minimize the loss?

How does the momentum work?

Momentum algorithm:

- Do a gradient descent step
- Apply the update from the last iteration, only smaller (momentum step)

$$w_{t+1} = w_t - \eta \nabla J(w_t) + \beta (w_t - w_{t-1})$$

Let's go to our tutorial to implement Momentum!

## How to minimize the loss?

Adaptive Methods <3

- Learning rate schedules

$$w_{t+1} = w_t - \eta_t \nabla J(w_t)$$

## How to minimize the loss?

Adaptive Methods <3

- Learning rate schedules

$$w_{t+1} = w_t - \eta_t \nabla J(w_t)$$

Polynomial schedules, e.g.  $\rightarrow \eta_t = \frac{\alpha}{c+t}, \eta_t = \frac{\alpha}{c+\sqrt{t}}$

Exponential

Stepwise decay

Cosine/cyclical schedules

## How to minimize the loss?

Adaptive Methods <3

- **Adagrad (Duchi et al., 2011)**
  - Adapts the learning rate for each parameter
  - Using running sum squares of past gradients
  - Typically used in stochastic form:
    - Instead of full gradient, use gradient from mini-batch

$$[w_{t+1}]_i = [w_t]_i - \frac{\eta}{\sqrt{[v_{t+1}]_i + \epsilon}} [\nabla J(w_t)]_i$$

$$[v_{t+1}]_i = \sum_{s=1}^t [\nabla J(w_s)]_i^2$$

## How to minimize the loss?

Adaptive Methods <3

- **RMSprop**
  - Uses a moving average instead of sum used by Adagrad
  - Moving average can be useful on non-convex objectives

$$[w_{t+1}]_i = [w_t]_i - \frac{\eta}{\sqrt{[v_{t+1}]_i + \epsilon}} [\nabla J(w_t)]_i$$

$$[v_{t+1}]_i = \alpha [v_t]_i + (1 - \alpha) [\nabla J(w_t)]_i^2$$

## How to minimize the loss?

Adaptive Methods <3

- **RMSprop**
  - Uses a moving average instead of sum used by Adagrad
  - Moving average can be useful on non-convex objectives

$$[w_{t+1}]_i = [w_t]_i - \frac{\eta}{\sqrt{[v_{t+1}]_i + \epsilon}} [\nabla J(w_t)]_i$$

$$[v_{t+1}]_i = \alpha[v_t]_i + (1 - \alpha)[\nabla J(w_t)]_i^2$$

- **Adam: RMSprop + momentum**

## How to minimize the loss?

Adaptive Methods <3

- **RMSprop**
  - Uses a moving average instead of sum used by Adagrad
  - Moving average can be useful on non-convex objectives

$$[w_{t+1}]_i = [w_t]_i - \frac{\eta}{\sqrt{[v_{t+1}]_i + \epsilon}} [\nabla J(w_t)]_i$$

$$[v_{t+1}]_i = \alpha [v_t]_i + (1 - \alpha) [\nabla J(w_t)]_i^2$$

- **Adam: RMSprop + momentum**

Let's go to our tutorial to implement RMSprop!

## How to minimize the loss?

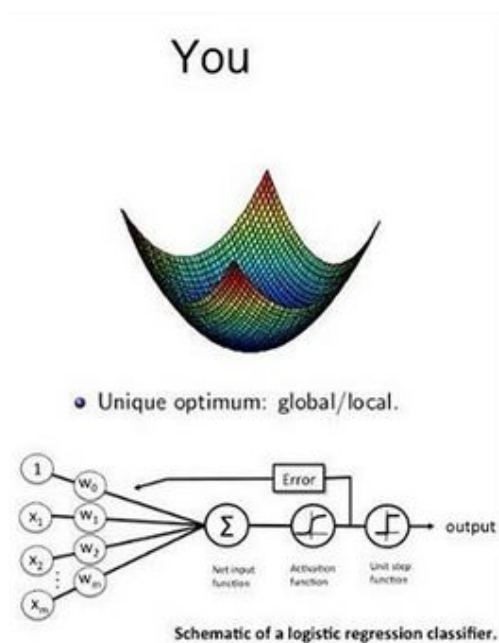
Non-convexity Problem



## How to minimize the loss?

### Non-convexity Problem

- Convex  $< 3$  -> have the same global and local minimum
- Non-convex -> have different global and local minimum

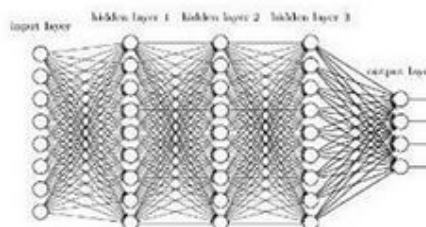


The guy she tells you  
not to worry about



- Multiple local optima
- In high dimensions possibly

Deep neural network



## How to minimize the loss?

### Non-convexity Problem

- Convex  $\rightarrow$  have the same global and local minimum
- Non-convex  $\rightarrow$  have different global and local minimum

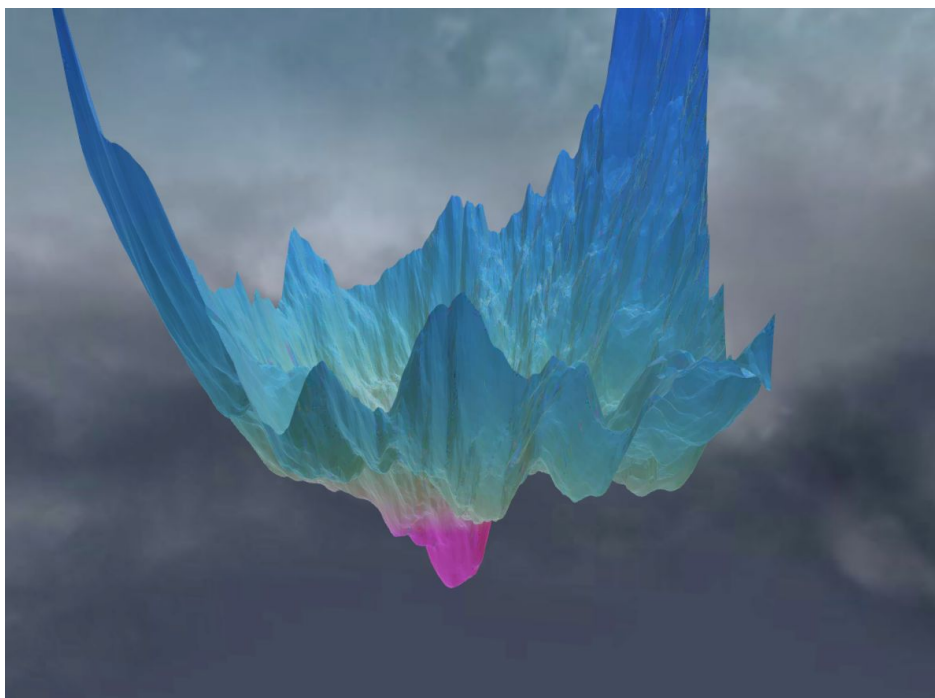
Great non-convexity, comes with great responsibility!

## How to minimize the loss?

Non-convexity Problem

- Initialization matters!!

<https://losslandscape.com/explorer>



## How to minimize the loss?

Non-convexity Problem

- Overparameterization to rescue!

**Let's go to tutorial and overparameterize our MLP!**

## How to minimize the loss?

Computation Cost Problem

## How to minimize the loss?

Computation Cost Problem

- Minibatch training  $< 3$

## How to minimize the loss?

Computation Cost Problem

- Minibatch training  $\rightarrow$  stochastic gradient descent

## How to minimize the loss?

### Computation Cost Problem

- Minibatch training  $\rightarrow$  stochastic gradient descent
  - Minibatch size:
    - Too small batch size: optimization bounces around a lot, and can lead to slower convergence to a minimum.
    - Too big batch size: won't fit on GPU
    - Rule of thumb  $\rightarrow$  pick the largest batch size that fits in the GPU



## How to minimize the loss?

### Computation Cost Problem

- Minibatch training  $\rightarrow$  stochastic gradient descent
  - Minibatch size:
    - Too small batch size: optimization bounces around a lot, and can lead to slower convergence to a minimum.
    - Too big batch size: won't fit on GPU
    - Rule of thumb  $\rightarrow$  pick the largest batch size that fits in the GPU

**Let's go to our tutorial and implement Mini Batch Training!!!**